

Baseline model (Logistic Regression):

	Set	Accuracy	Calibration
0	Train	0.680938	0.012042
1	Validation	0.682594	0.005245
2	Test	0.650000	0.047773

Algorithm 1- Information Theoretic Measures for Fairness-aware Feature Selection:

	Set	Accuracy	Calibration
0	Train	0.670462	0.036115
1	Validation	0.678043	0.022429
2	Test	0.636364	0.069751

Algorithm 2 – Learning Fair Representations:

```
Validation Accuracy: 0.5518  
Validation Calibration: 0.0673  
Test Accuracy: 0.5170  
Test Calibration: 0.1752
```

Accuracy: The baseline model is the most accurate across all three data splits (train, validation, test), with Algorithm 1 (FFS) following closely behind. Algorithm 2 (LFR) lags in accuracy on both validation and test sets.

Calibration: For calibration, a lower value indicates a fairer model. The baseline model has the best calibration on the test set, with Algorithm 1 (FFS) showing an increased calibration score, suggesting a decrease in fairness. Algorithm 2 (LFR) exhibits the highest calibration value on the test set, indicating further reduced fairness compared to the other models.